

# ENVT5503 time series in R

## Data Analysis in R for Environmental Science

### ENVT5503 Time Series

*Learning time series analysis concepts using hourly soil temperature data*

*Andrew Rate, School of Agriculture and Environment*



## Set up the R environment for time series analysis

We will need the additional functions in several R packages for specialized time series and other functions in R...

```
# Load the packages we need
library(zoo)          # for basic irregular time series functions
library(xts)          # we need the xts "eXtended Time Series" format for some functions
library(Kendall)      # for trend analysis with Mann-Kendall test
library(trend)        # for trend analysis using the Sen slope
library(forecast)     # for time series forecasting with ARIMA and exponential smoothing
library(tseries)      # for assessing stationarity using Augmented Dickey-Fuller test
library(lmtest)       # for Breusch-Pagan heteroscedasticity test etc.
library(car)          # for various commonly-used functions
library(ggplot2)      # alternative to base R plots

# (optional) make a better colour palette than the R default!
palette(c("black", "red3", "green3", "blue2",
          "darkcyan", "purple", "sienna", "gray50"))
par(mfrow = c(3,1), mar = c(4,4,1,1), mgp = c(1.7,0.3,0), tcl = 0.3, font.lab=2)
```

## Data input

Read the data into a data frame - this is how R often stores data - it's not the format we need but we'll use it for comparison.

## Non-time series object for comparison

```
soiltemp <- read.csv("soiltemp2.csv")
colnames(soiltemp) <- c("Date", "temp")
```

## Do some checks of the data

```
summary(soiltemp) # simple summary of each column
##      Date              temp
## Length:1464          Min.   :10.40
## Class :character      1st Qu.:12.45
## Mode  :character      Median :14.30
##                               Mean  :14.86
##                               3rd Qu.:17.41
##                               Max.   :21.70
str(soiltemp) # more detailed information about the R object ('str'=structure)
## 'data.frame':    1464 obs. of  2 variables:
## $ Date: chr  "2014-04-01 00:00:00" "2014-04-01 01:00:00" "2014-04-01 02:00:00" "2014-
04-01 03:00:00" ...
## $ temp: num  20.1 19.8 19.6 19.2 19 ...
```

## Check with a plot

```
plot(soiltemp$temp, type = "l", col = 6)
```




Figure 1: Plot of soil temperature time series data which is not (yet) formatted as a time series.

The horizontal axis in Figure 1 is just the row number of the data frame, not a date or time. We really want the data in a different type of R object! We use the `read.csv.zoo()` function in the `zoo` package to read the data in as a time series. This gives a data object of class 'zoo' which is much more flexible than the default time series object in R.

```
soiltemp_T15_zoo <- read.csv.zoo("soiltemp2.csv",
                                format = "%Y-%m-%d %H:%M:%S",
                                tz = "Australia/Perth",
                                index.column=1,
                                header = TRUE)

# do some quick checks of the new R object:
summary(soiltemp_T15_zoo)
##      Index              soiltemp_T15_zoo
## Min.   :2014-04-01 00:00:00   Min.   :10.40
## 1st Qu.:2014-04-16 05:45:00   1st Qu.:12.45
## Median :2014-05-01 11:30:00   Median :14.30
## Mean   :2014-05-01 11:30:00   Mean   :14.86
## 3rd Qu.:2014-05-16 17:15:00   3rd Qu.:17.41
## Max.   :2014-05-31 23:00:00   Max.   :21.70
str(soiltemp_T15_zoo) # POSIXct in the output is a date-time format
## 'zoo' series from 2014-04-01 to 2014-05-31 23:00:00
##  Data: num [1:1464] 20.1 19.8 19.6 19.2 19 ...
##  Index: POSIXct[1:1464], format: "2014-04-01 00:00:00" "2014-04-01 01:00:00" "2014-04-01 02:00:00" ...
```

It's usually useful to check our data with a plot (Figure 2)

```
plot(soiltemp_T15_zoo, col = 2)
```

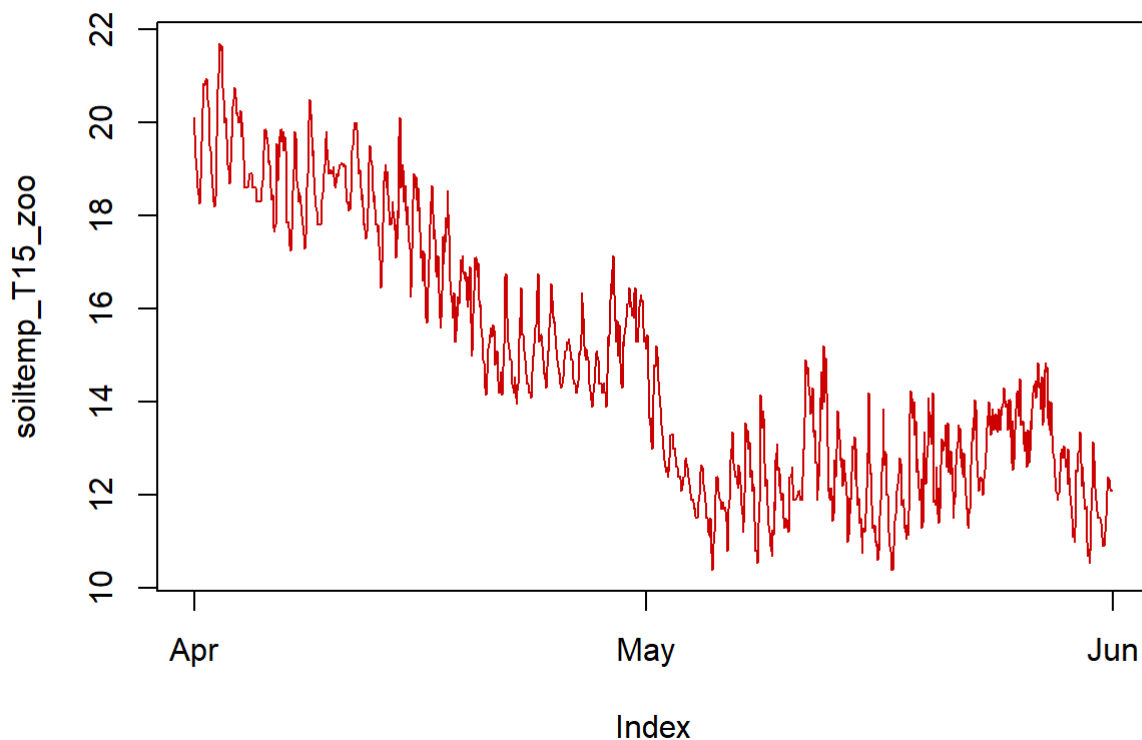


Figure 2: Plot of soil temperature time series data which is formatted as a zoo time series object.

Sometimes we need to use another time series data format, `xts` (eXtended Time Series) which allows us to use more functions...

Make an xts object from our zoo object and check:

```
soiltemp_T15_xts <- as.xts(soiltemp_T15_zoo)
str(soiltemp_T15_xts) # just to check
## An 'xts' object on 2014-04-01/2014-05-31 23:00:00 containing:
##   Data: num [1:1464, 1] 20.1 19.8 19.6 19.2 19 ...
##   Indexed by objects of class: [POSIXct,POSIXt] TZ: Australia/Perth
##   xts Attributes:
##    NULL
plot(soiltemp_T15_xts, col = 6, ylab = "Soil temperature (\u00B0C)") # just to check
```




Figure 3: Plot of soil temperature time series data which is formatted as an xts time series object.

The plot in Figure 3 shows that plotting an `xts` time series object gives a somewhat more detailed plot than for a `zoo` formatted time series object (Figure 2).

## Exploratory data analysis of time series

We first examine a plot of the data in our time series object, both the raw data and using common transformations (Figure 4). We compare with transformed time series data - we may need to do this to meet later modelling assumptions.

First we change the default plotting parameters using `par(...)`, then plot the object and some transformed versions (with custom axis labels).

```
par(mfrow = c(3,1), mar = c(4,4,1,1), mgp = c(1.7,0.3,0), tcl = 0.3, font.lab=2)

plot(soiltemp_T15_zoo, ylab = "Temperature (\u00B0C)",
     xlab = "Date", col = 7, lwd = 2, cex.lab = 1.4)

plot(log10(soiltemp_T15_zoo),
     ylab = expression(bold(paste(log[10], "(Temperature, \u00B0C)"))),
     xlab = "Date", col = 6, lwd = 2, cex.lab = 1.4)

pt0 <- powerTransform(coredata(soiltemp_T15_zoo))
if(pt0$lambda<0) {
  plot(
    -1 * (soiltemp_T15_zoo ^ pt0$lambda),
    ylab = "power-transf. (Temp., \u00B0C)",
    xlab = "Date",
    col = 4, lwd = 2, cex.lab = 1.4
  )
} else {
  plot((soiltemp_T15_zoo ^ pt0$lambda),
       ylab = "power-transf. (Temp., \u00B0C)",
```

```

xlab = "Date",
col = 4, lwd = 2, cex.lab = 1.4
)
}

```




Figure 4: Plots of soil temperature time series data as untransformed, log10-transformed, and power transformed values.

Which of these plots looks like it might be homoscedastic, *i.e.*, have constant variance regardless of time?

The R Cookbook suggests Box-Cox (power) transforming the variable to “stabilize the variance” (*i.e.* reduce heteroscedasticity). This does work, but can be better to use  $\log_{10}[\text{conc}]$ , as the values are easier to interpret.

If we think a transformation is needed, then run something like the code below (which does a square root transformation, *i.e.*  $\text{variable}^{0.5}$ ). This is just an example – we may, for example, decide that the  $\log_{10}$ -transformation is more suitable.

```

soiltemp_T15_zoo <- soiltemp_T15_zoo^0.5
# don't forget the xts version either!
soiltemp_T15_xts <- as.xts(soiltemp_T15_zoo)

```

## Assessing if a time series variable is stationary

A **stationary** variable's mean and variance are not dependent on time. In other words, for a stationary series, the mean and variance at any time are representative of the whole series.

If there is a trend for the value of the variable to increase or decrease, or if there are periodic fluctuations, we don't have a stationary time series.

Many useful statistical analyses and models for time series models need a stationary time series as input, or a time series that can be made stationary with transformations or differencing.

## Testing for stationarity

```

# we need the package 'tseries' for the Augmented Dickey-Fuller (adf) Test
d0 <- adf.test(soiltemp_T15_zoo); print(d0)
##
## Augmented Dickey-Fuller Test
##
## data: soiltemp_T15_zoo
## Dickey-Fuller = -3.9515, Lag order = 11, p-value = 0.01141
## alternative hypothesis: stationary
d1 <- adf.test(diff(soiltemp_T15_zoo,1)); print(d1)
## Warning in adf.test(diff(soiltemp_T15_zoo, 1)): p-value smaller than printed p-
## value
##
## Augmented Dickey-Fuller Test
##
## data: diff(soiltemp_T15_zoo, 1)
## Dickey-Fuller = -16.143, Lag order = 11, p-value = 0.01
## alternative hypothesis: stationary
d2 <- adf.test(diff(diff(soiltemp_T15_zoo,1),1)); print(d2)
## Warning in adf.test(diff(diff(soiltemp_T15_zoo, 1), 1)): p-value smaller than
## printed p-value
##
## Augmented Dickey-Fuller Test
##
## data: diff(diff(soiltemp_T15_zoo, 1), 1)
## Dickey-Fuller = -17.204, Lag order = 11, p-value = 0.01
## alternative hypothesis: stationary

```

## Plot differencing for stationarity

The plots in Figure 5 below show that constant mean and variance (*i.e.* stationarity) is achieved after a single differencing step. The `adf.test` p-value suggests that the non-differenced series is already stationary, but the first sub-plot in Figure 5 does not support this.

```

par(mfrow = c(3,1), mar = c(0,4,0,1), oma = c(4,0,1,0), cex.lab = 1.4,
    mgp = c(2.5,0.7,0), font.lab = 2)
plot(soiltemp_T15_zoo, ylab = "Raw data, no differencing",
     xlab="", xaxt="n", col = 8)
lines(loess.smooth(index(soiltemp_T15_zoo),coredata(soiltemp_T15_zoo)),
     col = 4, lwd = 2)
abline(lm(coredata(soiltemp_T15_zoo) ~ index(soiltemp_T15_zoo)), col = 2)
legend("topright", legend = c("soiltemp_T15_Data", "Loess smoothing", "Linear model"),
     cex = 1.8, col = c(1,4,2), lwd = c(1,2,1), bty = "n")
mtext(paste("adf.test p value =", signif(d0$p.value,3)),
     side = 1, line = -1.2, adj = 0.05)
plot(diff(soiltemp_T15_zoo,1),
     ylab = "First differencing",
     xlab="", xaxt="n", col = 8)
abline(h = 0, col = "grey", lty = 2)
lines(loess.smooth(index(diff(soiltemp_T15_zoo,1)),coredata(diff(soiltemp_T15_zoo,1))),
     col = 4, lwd = 2)

```

```

abline(lm(coredata(diff(soiltemp_T15_zoo,1)) ~
            index(diff(soiltemp_T15_zoo,1))), col = 2)
mtext(paste("adf.test p value =", signif(d1$p.value,3)),
      side = 1, line = -1.2, adj = 0.05)
plot(diff(diff(soiltemp_T15_zoo,1),1),
      ylab = "Second differencing",
      xlab="Date", col = 8)
mtext("Date",side = 1, line = 2.2, font = 2)
abline(h = 0, col = "grey", lty = 2)
lines(loess.smooth(index(diff(diff(soiltemp_T15_zoo,1),1)),
                    coredata(diff(diff(soiltemp_T15_zoo,1),1))),
      col = 4, lwd = 2)
abline(lm(coredata(diff(diff(diff(soiltemp_T15_zoo,1))) ~
            index(diff(diff(diff(soiltemp_T15_zoo,1))))), col = 2)
mtext(paste("adf.test p value =", signif(d2$p.value,3)),
      side = 1, line = -1.2, adj = 0.05)

```




Figure 5: Plots of soil temperature time series data and its first and second differences, with each sub-plot showing the smoothed (Loess) and linear trends.

## Finding the trend

### 1. Determine if there is actually any trend

(see Department of Water 2015).

#### 1a. Apply the Mann-Kendall test from the 'trend' package

```
mk.test(coredata(soiltemp_T15_zoo))
##
## Mann-Kendall trend test
##
## data: coredata(soiltemp_T15_zoo)
## z = -36.387, n = 1464, p-value < 2.2e-16
## alternative hypothesis: true S is not equal to 0
## sample estimates:
##          S          varS          tau
## -6.797270e+05  3.489630e+08 -6.371125e-01
## or
## SeasonalMannKendall(soiltemp_T15_zoo) # needs base R regular time series object
```

The output from `mk.test()` shows a negative slope (*i.e.* the value of `S`), with the `p-value` definitely < 0.05, so we can reject the null hypothesis of zero slope.

#### 1b. Estimate the Sen's slope

```
sens.slope(coredata(soiltemp_T15_zoo))
##
## Sen's slope
##
## data: coredata(soiltemp_T15_zoo)
## z = -36.387, n = 1464, p-value < 2.2e-16
## alternative hypothesis: true z is not equal to 0
## 95 percent confidence interval:
## -0.005963303 -0.005570410
## sample estimates:
## Sen's slope
## -0.005765595
```

The output from `sens.slope()` also shows a negative slope (*i.e.* the value of `Sen's slope` at the end of the output block). Again, the `p-value` is < 0.05, so we can reject the null hypothesis of zero slope. The `95 percent confidence interval` does not include zero, also showing that the Sen's slope is significantly negative in this example.

## 2. Visualising a trend using a moving average

We create a new time series from a moving average for each 24h to remove daily periodicity using the `rollmean()` function from the `zoo` package. We know these are hourly data with diurnal fluctuation, so a rolling mean for chunks of length 24 should be OK, and this seems to be true based on Figure 6. In some cases the `findfrequency()` function from the `xts` package can detect the periodic frequency for us. A moving average will always smooth our data, so a smooth curve doesn't necessarily mean that we have periodicity (seasonality).

```
ff <- findfrequency(soiltemp_T15_xts) # often an approximation!! check the data!
cat("Estimated frequency is",ff,"\n")
## Estimated frequency is 24
soiltemp_T15_movAv <- rollmean(soiltemp_T15_zoo, 24)
plot(soiltemp_T15_zoo, col=8, type="l") # original data
# add the moving average
lines(soiltemp_T15_movAv, lwd = 2)
legend("topright", legend = c("Data","Moving average"), col = c(8,1), lwd = c(1,2))
```






Figure 6: Plots of soil temperature time series data and the 24-hour moving average.

### 3. Showing a trend using a linear (regression) model

First we create a linear model of the time series...

```
lm0 <- lm(coredata(soiltemp_T15_zoo) ~ index(soiltemp_T15_zoo))
summary(lm0)
##
## Call:
## lm(formula = coredata(soiltemp_T15_zoo) ~ index(soiltemp_T15_zoo))
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -3.9217 -1.0968  0.0743  1.0938  3.5318
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2.213e+03  3.370e+01   65.66  <2e-16 ***
## index(soiltemp_T15_zoo) -1.571e-06  2.409e-08  -65.22  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.402 on 1462 degrees of freedom
## Multiple R-squared:  0.7442, Adjusted R-squared:  0.744
## F-statistic: 4253 on 1 and 1462 DF, p-value: < 2.2e-16
```

...then use a plot (Figure 7) to look at the linear relationship.

```
soiltemp_T15_lmfit <- zoo(lm0$fitted.values, index(soiltemp_T15_zoo))
plot(soiltemp_T15_zoo, col = 8, type = "l")
lines(soiltemp_T15_lmfit, col = 2, lty = 2)
legend("topright", legend = c("Data", "Linear trend"), col = c(8, 2), lty = c(1, 2))
```

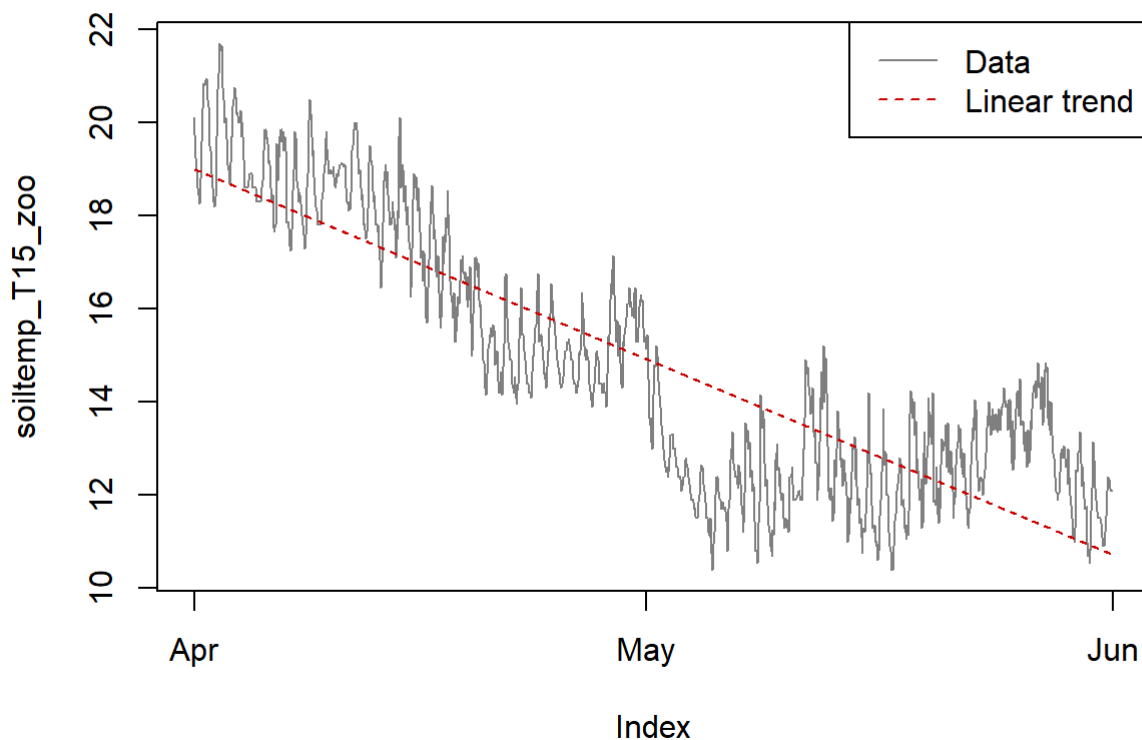


Figure 7: Plots of soil temperature time series data and the overall trend shown by a linear model.

Figure 7 shows that a linear model does not really capture the trend (ignoring the periodicity) very convincingly. The next option 'loess' smoothing option below tries to present the trend more accurately.

#### 4. Using Locally Estimated Scatterplot Smoothing (loess)

(**loess**, sometimes called 'lowess' is a form of locally weighted non-parametric regression to fit smooth curves to data). The amount of smoothing is controlled by the `span =` parameter in the `loess.smooth()` function (from base R).

**NOTE:** loess smoothing does not have the relationship to ARIMA that moving averaging does, but does allow us to separate periodicity from random error in time series decomposition.

```
y_trend <- loess.smooth(index(soiltemp_T15_zoo),
                        coredata(soiltemp_T15_zoo),
                        span = 0.15, evaluation = length(soiltemp_T15_zoo))
plot(soiltemp_T15_zoo, col = 8, type = "l")
soiltemp_T15_trend <- zoo(y_trend$y, index(soiltemp_T15_zoo))
lines(soiltemp_T15_lmfit, col = 2, lty = 2)
lines(soiltemp_T15_trend, col = "skyblue", lwd = 3)
legend("topright", bty = "n", inset = 0.02, cex = 1.25,
      legend = c("actual data", "linear model", "loess smoothed"),
      col = c(8, 2, "skyblue"), lty=c(1, 2, 1), lwd = c(1, 1, 3))
```

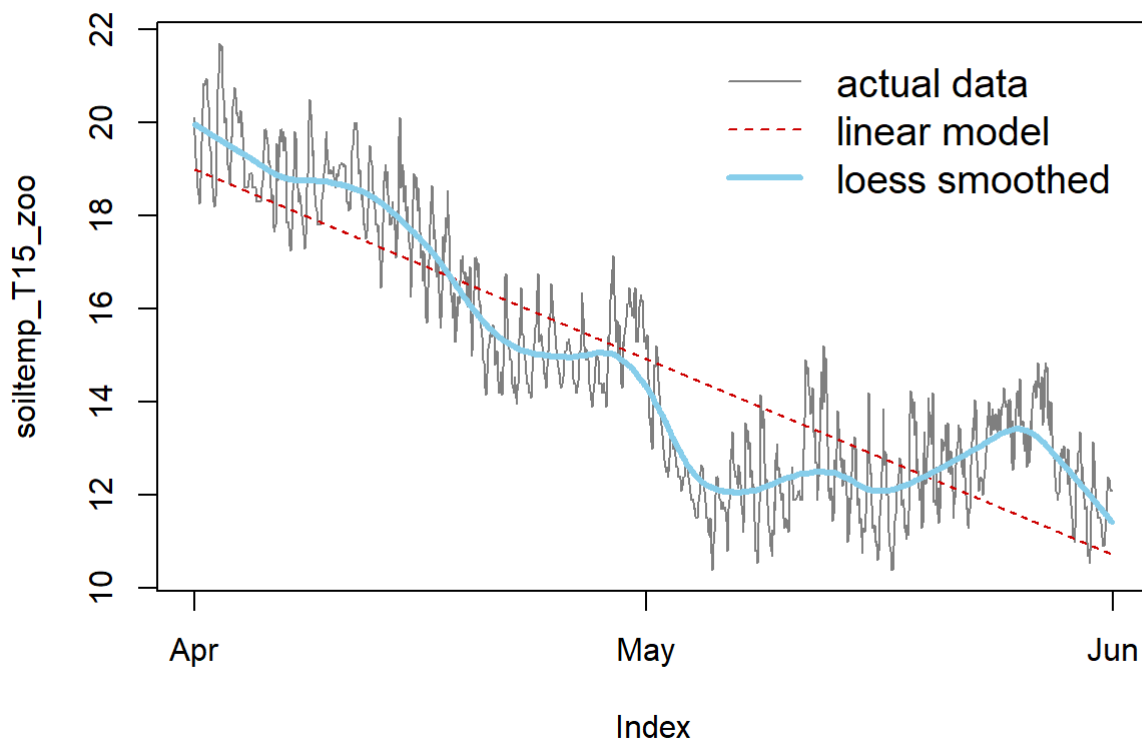


Figure 8: Plots of soil temperature time series data and the overall trend shown by a loess smoothing model.

## Isolating the Time Series Periodicity

**NOTE THAT TIME SERIES DON'T ALWAYS HAVE PERIODICITY !**

To model the periodicity we need to understand the autocorrelation structure of the time series. We can do this graphically, first by plotting the autocorrelation function `acf()` which outputs a plot by default:

```
acf(soiltemp_T15_xts, main = "")
```




Figure 9: The autocorrelation function plot for soil temperature time series data (the horizontal axis is in units of seconds).

What does Figure 9 tell you about autocorrelation in this time series?

Then plot the partial autocorrelation function (`pacf()`)

```
pacf(soiltemp_T15_xts, main = "")
```




Figure 10: The partial autocorrelation function plot for soil temperature time series data (horizontal axis units are seconds).

Interpreting partial autocorrelations (Figure 10) is more complicated - refer to Long & Teetor (2019, Section 14.15). Simplistically, partial autocorrelation allows us to identify which and how many autocorrelations will be needed to model the time series data.

### We use the 'Box-Pierce test' for autocorrelation

The null hypothesis  $H_0$  is that no autocorrelation exists at any lag distance (so  $p \leq 0.05$  'rejects'  $H_0$ ):

```
Box.test(soiltemp_T15_xts)
##
## Box-Pierce test
##
## data: soiltemp_T15_xts
## X-squared = 1436.3, df = 1, p-value < 2.2e-16
Box.test(diff(soiltemp_T15_xts,1)) # 1 difference
##
## Box-Pierce test
##
## data: diff(soiltemp_T15_xts, 1)
## X-squared = 15.658, df = 1, p-value = 7.589e-05
Box.test(diff(diff(soiltemp_T15_xts,1),1)) # 2 differences
##
## Box-Pierce test
##
## data: diff(diff(soiltemp_T15_xts, 1), 1)
## X-squared = 357.45, df = 1, p-value < 2.2e-16
```

## Make a time series of the loess residuals

Remember that we made a loess model of the time series ... the residuals (Figure 11) can give us the combination of periodicity component plus any random variation.

```
soiltemp_T15_periodic <- soiltemp_T15_zoo - soiltemp_T15_trend
plot(soiltemp_T15_trend, ylim = c(-2,20), lty = 3)
lines(soiltemp_T15_periodic, col = "coral", lwd = 2) # just to check
legend("topright", legend = c("LOESS trend", "Periodicity plus noise"),
      col = c(1,"coral"), lwd = c(1,2), lty = c(3,1))
```

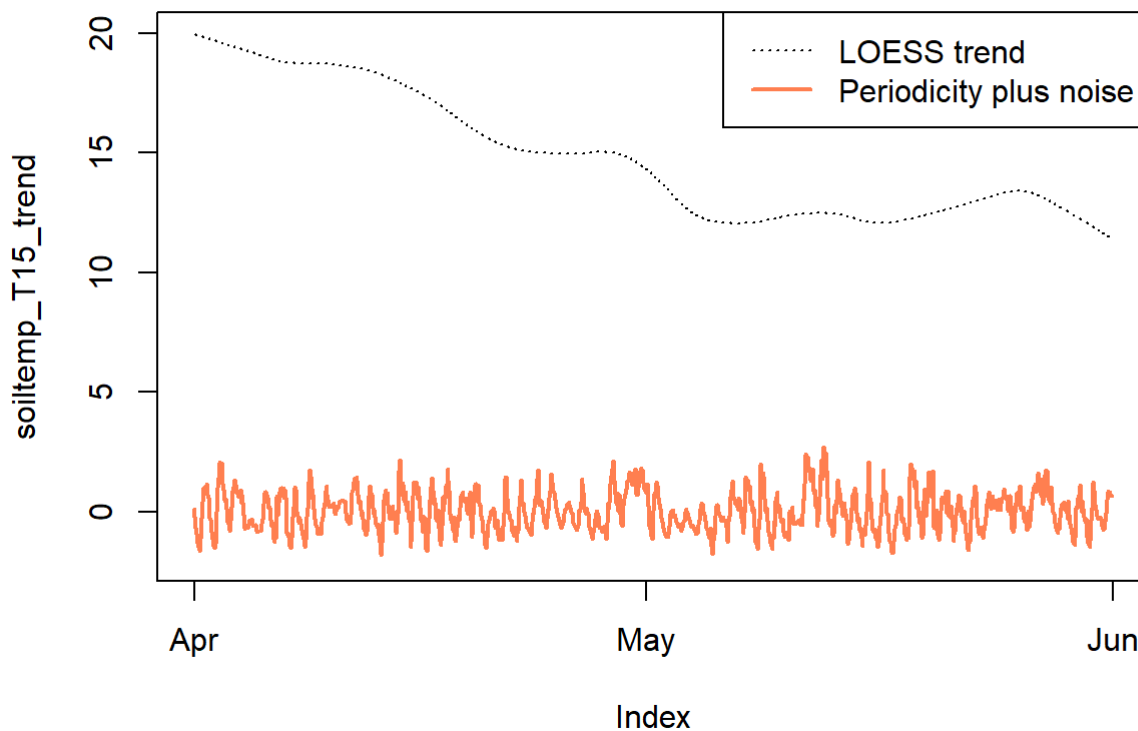


Figure 11: An incomplete decomposition of the soil temperature time series into smoothed trend and periodicity plus error (noise) components.

We can also use *less smoothing* in the loess function to *retain* periodicity; we adjust the `span` option (lower values of `span` give less smoothing; we just need to experiment with different values). The difference between the data and the less-smoothed loess should be just 'noise' or 'error'.

We first generate a loess model which is stored in the object `temp_LOESS2`:

```
temp_LOESS2 <- loess.smooth(index(soiltemp_T15_zoo),  
                             coredata(soiltemp_T15_zoo),  
                             span = 0.012, evaluation = length(soiltemp_T15_zoo))
```

We then use the new loess model to make a time series which contains both periodic and trend information:

```
soiltemp_T15_LOESS2 <- zoo(temp_LOESS2$y, index(soiltemp_T15_zoo))
```

The difference between the data and the less-smoothed loess should be just 'noise' or 'error', so we make a new time series based on this difference:

```
soiltemp_T15_err <- soiltemp_T15_zoo - soiltemp_T15_LOESS2
```

plot, setting y axis limits to similar scale to original data:

```
par(mar = c(4,4,1,1), mgp = c(1.7,0.3,0), tcl = 0.3, font.lab = 2)  
plot(soiltemp_T15_trend, ylim = c(-2,20), lty = 2, lwd = 2, # from above  
      xlab = "2014 Date", ylab = "Soil temperature (\u00B0C)")  
lines(soiltemp_T15_LOESS2, col = 3, lwd = 2)  
lines(soiltemp_T15_err, col = 6) # from a couple of lines above  
legend("left",  
       legend = c("Trend (coarse LOESS)",  
                  "Periodicity + trend (fine LOESS)",  
                  "Unexplained variation"),  
       col = c(1,3,6), lwd = c(2,2,1), lty = c(2,1,1), bty="n")
```

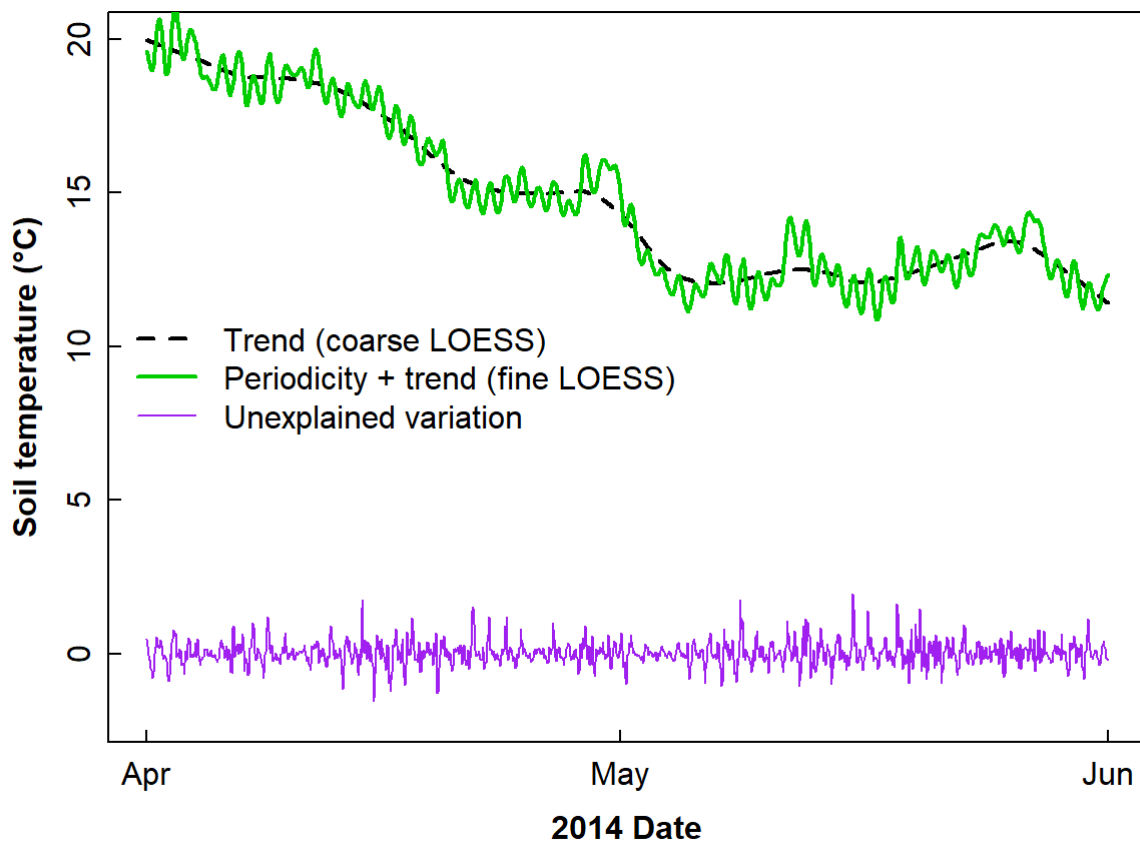


Figure 12: Incomplete decomposition of soil temperature time series data showing the smoothed *trend*, combined *trend + periodicity*, and *error* components.

The periodicity by itself should be represented by the difference between the very smoothed (trend) and less smoothed (trend + periodicity) loess, as suggested by the two loess curves in Figure 12. We now make yet another time series object to hold this isolated periodicity:

```
soiltemp_T15_periodic2 <- soiltemp_T15_LOESS2 - soiltemp_T15_trend
```

### Plot everything to show the components of time series decomposition

```
plot(cbind(soiltemp_T15_zoo, soiltemp_T15_periodic2,
           soiltemp_T15_trend, soiltemp_T15_err), main = "", xlab = "Date in 2014",
     cex.main = 1.5, yax.flip = TRUE, col = c(8,6,4,2), ylim = c(-2,22))
x1 <- (0.5*(par("usr")[2]-par("usr")[1]))+par("usr")[1]
y <- (c(0.2, 0.3, 0.6, 0.82)*(par("usr")[4]-par("usr")[3]))+par("usr")[3]
text(rep(x1,4), y,
     labels = c("Unaccounted variation", "Trend", "Apparent periodicity", "Data"),
     col = c(2,4,6,8))
```

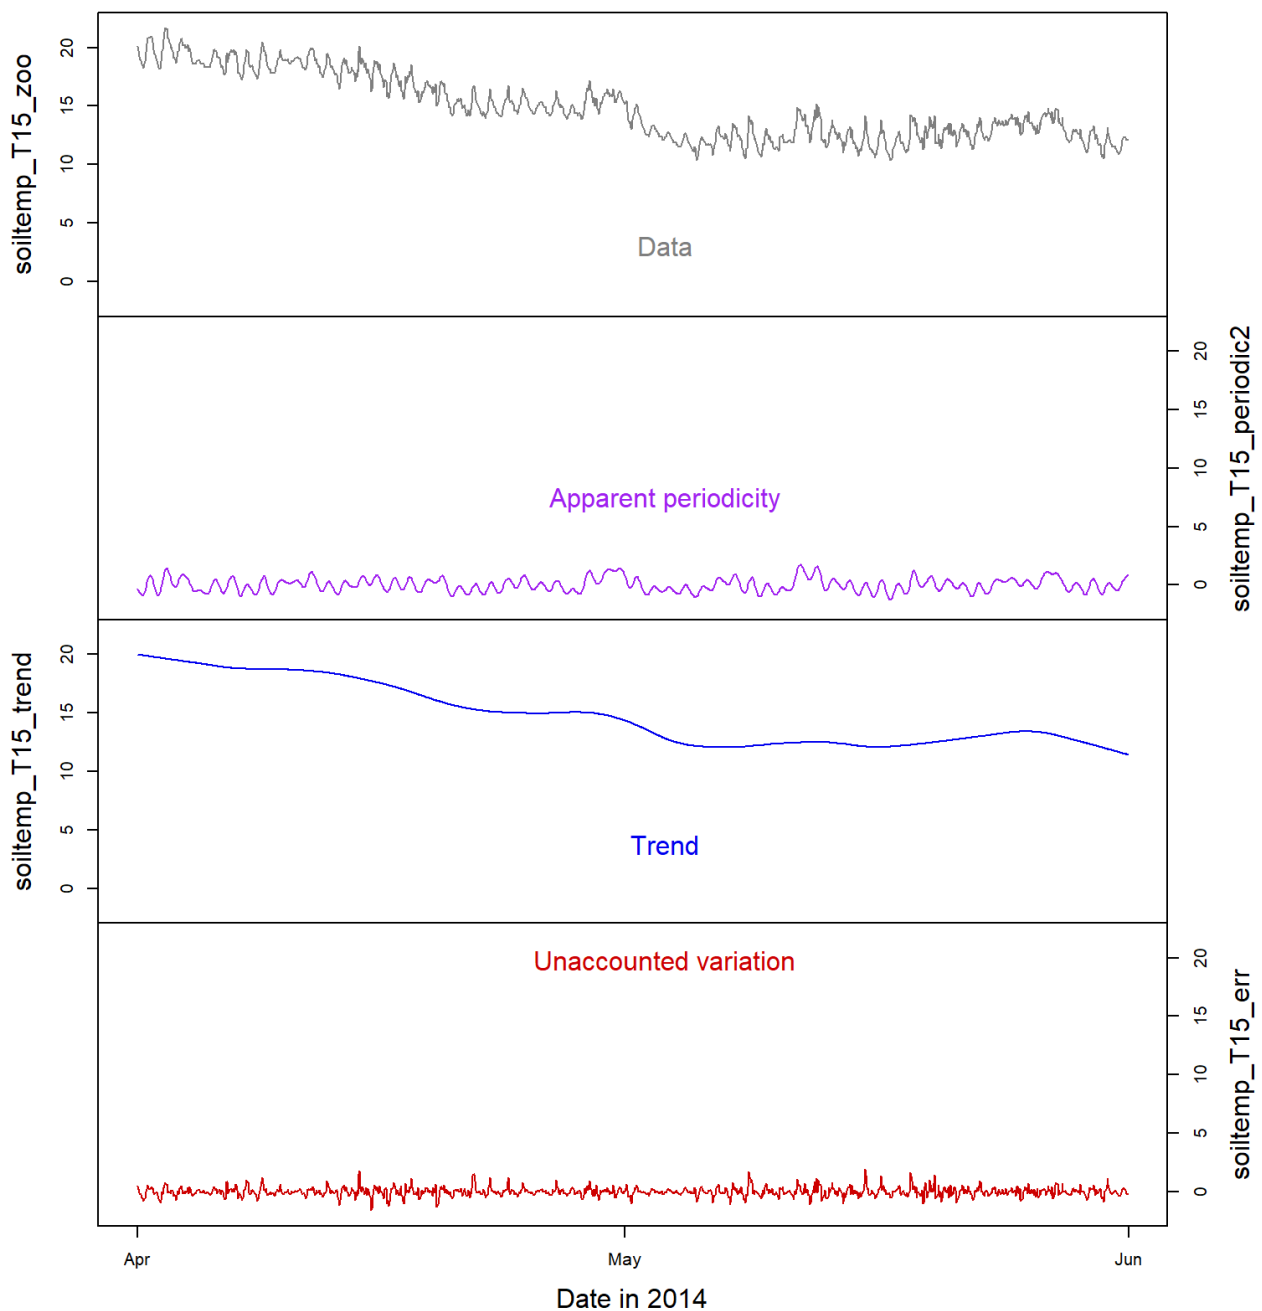


Figure 13: Time series decomposition into the three main components (periodicity, trend, and error) for soil temperature ( $^{\circ}\text{C}$ ) at 15 cm depth.

Figure 13 shows the original time series data, and the three important components of the decomposed time series: the *periodicity*, the *trend*, and the *error* or unexplained variation. We should remember that:

- this time series was relatively easy to decompose into its components, but not all time series will be so ‘well-behaved’;
- there are different ways in which we can describe the trend and periodicity (we have just used the loess smoothing functions for clear visualization).

**This ends our exploratory data analysis of time series, which leads us into ARIMA forecast modelling.**



# Modelling time series with ARIMA

All the analysis of our data before ARIMA is really exploratory data analysis of time series:

- Does our time series have periodicity?
- Can we get stationarity with a moving average?
- Does our time series have autocorrelation?
- Can we get stationarity by differencing?

All of these operations are possible components of **ARIMA** models!

We recommend using the xts format of a time series in ARIMA model functions and forecasting.

Use the **forecast** R package to run an ARIMA model

```
auto.arima(soiltemp_T15_xts, max.p = 3, max.q = 3, max.d = 0,
            seasonal = TRUE)
## Series: soiltemp_T15_xts
## ARIMA(1,0,3) with non-zero mean
##
## Coefficients:
##          ar1      ma1      ma2      ma3      mean
##          0.9883  0.0923  0.0892  0.0999  15.0050
## s.e.      0.0041  0.0263  0.0255  0.0258  0.9301
##
## sigma^2 = 0.1188: log likelihood = -517.87
## AIC=1047.73   AICc=1047.79   BIC=1079.47
```

The `auto.arima()` function runs a complex algorithm (Hyndman *et al.* 2020) to automatically select the best ARIMA model on the basis of the **Aikake Information Criterion** (AIC), a statistic which estimates how much information is 'lost' in the model compared with reality. The AIC combines how well the model describes the data with a 'penalty' for increased complexity of the model. Using AIC, a better fitting model might not be selected if it has too many predictors. The best ARIMA models will have the **lowest** AIC (or AICc) value.

Use output of **auto.arima()** to run arima

The `auto-arima()` algorithm is not perfect! – but it does provide a starting point for examining ARIMA models.

The output of the `auto-arima()` function includes a description of the best model the algorithm found, shown as **ARIMA(p,d,q)**. The p,d,q refer to

Parameter	Meaning	Informed by
p	The number of autoregressive predictors	Partial autocorrelation
d	The number of differencing steps	Stationarity tests ± differencing
q	The number of moving average predictors	Stationarity tests ± moving averages

A periodic or *seasonal* ARIMA model (often called SARIMA) has a more complex specification:

**ARIMA(p,d,q) (P,D,Q) (n)** ,

where the additional parameters refer to the seasonality: **P** is the number of seasonal autoregressive predictors, **D** the seasonal differencing, **Q** the seasonal moving averages, and **n** is the number of time steps per period/season.

## 1. With no seasonality

```
am0 <- arima(x = soiltemp_T15_xts, order = c(1,0,3))
summary(am0)
confint(am0)
##
```

```
## Call:
## arima(x = soiltemp_T15_xts, order = c(1, 0, 3))
##
## Coefficients:
##          ar1          ma1          ma2          ma3  intercept
##          0.9883  0.0923  0.0892  0.0999    15.0050
## s.e.    0.0041  0.0263  0.0255  0.0258     0.9301
##
## sigma^2 estimated as 0.1184:  log likelihood = -517.87,  aic = 1047.73
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE          MASE
## Training set -0.005215705 0.3441525 0.2467179 -0.08670514 1.716531 0.980305
##              ACF1
## Training set -0.0001019764
##              2.5 %          97.5 %
## ar1          0.98017359  0.9963265
## ma1          0.04074353  0.1438382
## ma2          0.03916128  0.1392272
## ma3          0.04923344  0.1504821
## intercept 13.18199812 16.8279311
```

The output from `summary(am0)` shows the values and uncertainties of the model parameters (`ar1`, `ma1`, `ma2`, `ma3`, `intercept`), the goodness-of-fit parameters (we're interested in `aic`).

**Note:** `ar` parameters are auto-regression coefficients, and `ma` are moving average coefficients.

The output from `confint(am0)` is a table showing the 95% confidence interval for the parameters. The confidence intervals should not include zero! (if so, this would mean we can't be sure if the parameter is useful or not).

## 2. With seasonality

```
ff <- findfrequency(soiltemp_T15_xts)
cat("Estimated time series frequency is",ff,"\n")
aml <- arima(x = soiltemp_T15_xts, order = c(1,0,2),
             seasonal = list(order = c(1, 0, 1), period = ff))
summary(aml)
confint(aml)
## Estimated time series frequency is 24
##
## Call:
## arima(x = soiltemp_T15_xts, order = c(1, 0, 2), seasonal = list(order = c(1,
## 0, 1), period = ff))
##
## Coefficients:
##          ar1          ma1          ma2          sar1          smal  intercept
##          0.9968  -0.1466  -0.0984  0.9998  -0.9856    14.829
## s.e.    0.0020  0.0268  0.0288  0.0004  0.0143    10.415
##
## sigma^2 estimated as 0.09073:  log likelihood = -355.68,  aic = 725.35
##
## Training set error measures:
##              ME          RMSE          MAE          MPE          MAPE          MASE
## Training set -0.01572894 0.30122 0.2169846 -0.1427709 1.504911 0.8621632
##              ACF1
## Training set 0.0006582706
##              2.5 %          97.5 %
## ar1          0.9928401  1.00080710
## ma1          -0.1990623 -0.09412079
## ma2          -0.1549117 -0.04197365
## sar1          0.9989639  1.00061967
## smal         -1.0136981 -0.95746903
## intercept -5.5840446 35.24194848
```

**Note:** `sar` parameters are seasonal/ auto-regression coefficients, and `sma` are seasonal/ moving average coefficients.

The model which includes periodicity has the lowest AIC value (which is not surprising, since the soil temperature data have clear periodicity).

Checking residuals using the `checkresiduals()` function from the `forecast` package is our best diagnostic tool to assess the validity of our models. [This is a separate issue from how well the model describes the data, measured by AIC.]

In the output (plot and text) from `checkresiduals()`: - the residual plot (top) should look like white noise - the residuals should not be autocorrelated (bottom left plot) - the p-value from the Ljung-Box test should be  $> 0.05$  (text output) - the residuals should be normally distributed (bottom right plot)

```
checkresiduals(am0)
```

Residuals from ARIMA(1,0,3) with non-zero mean

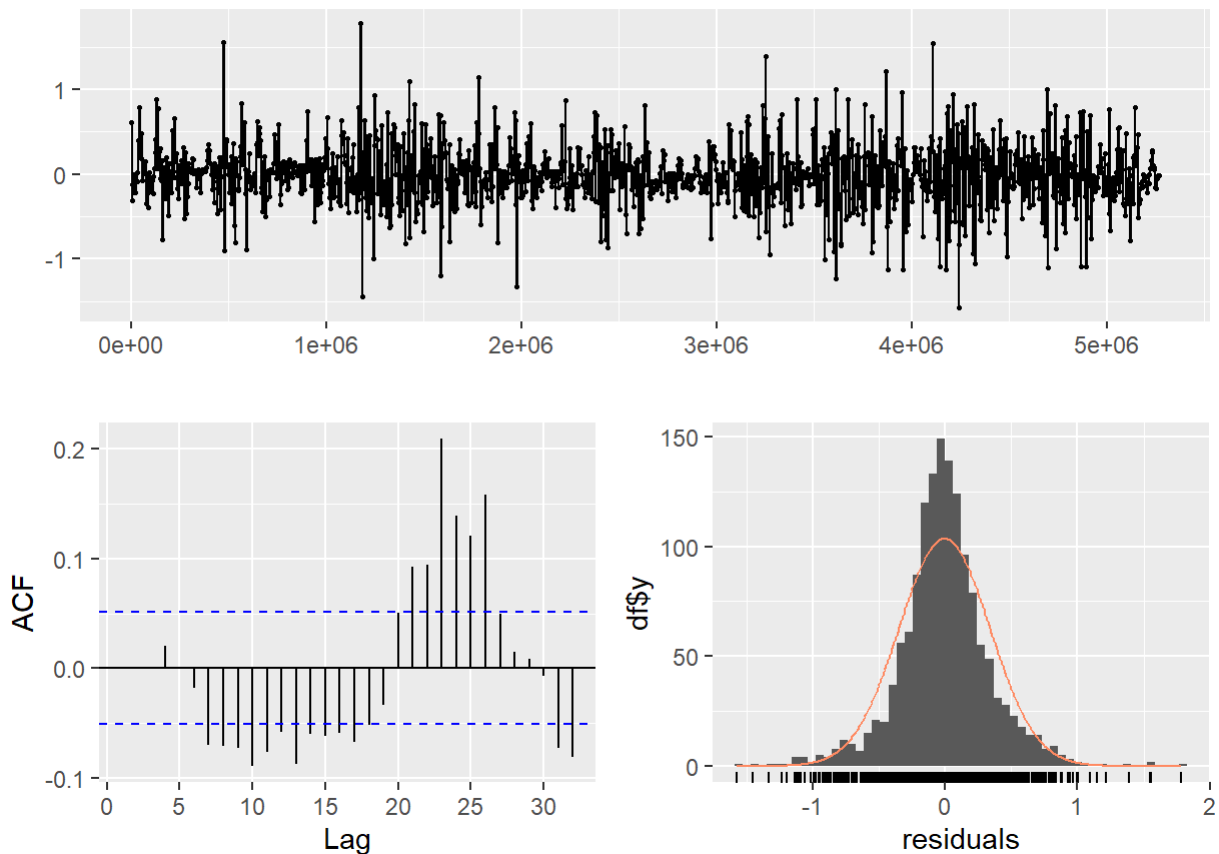


Figure 14: Residual diagnostic plots for a non-seasonal ARIMA model of the soil temperature time series.

```
##  
##  Ljung-Box test  
##  
## data:  Residuals from ARIMA(1,0,3) with non-zero mean  
## Q* = 35.092, df = 6, p-value = 4.136e-06  
##  
## Model df: 4.    Total lags used: 10  
checkresiduals(am1)
```




Figure 15: Residual diagnostic plots for a seasonal ARIMA (SARIMA) model of the soil temperature time series.

```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(1,0,2)(1,0,1)[24] with non-zero mean
## Q* = 20.04, df = 5, p-value = 0.001228
##
## Model df: 5.   Total lags used: 10
```

In both the diagnostic plots (Figure 14 and Figure 15) the residuals appear to be random and normally distributed. There seems to be a more obvious autocorrelation of residuals for the non-seasonal model, reinforcing the conclusion from the AIC value that the **seasonal** ARIMA model is the most appropriate.

## Use the ARIMA model to produce a forecast using both models

The whole point of generating ARIMA models is so that we can attempt to forecast the future trajectory of our time series based on our existing time series data.

To do this, we use the appropriately named `forecast()` function from the `forecast` package. The option `h` is to specify how long we want to forecast for after the end of our data – in this case 168 hours (same units as the periodicity from `findfrequency()`), that is, one week.

```
fc0 <- forecast(am0, h = 168)
fc1 <- forecast(am1, h = 168)
```

## Then, look at forecasts with plots

```
par(mfrow = c(2, 1), cex.main = 0.9, mar = c(0,3,0,1), oma = c(3,0,1,0),
    mgp=c(1.6,0.3,0), tcl=0.2, font.lab=2)
plot(fc0, ylab = "Temperature (\u00B0C)",
     fcol = 4, xlab="", main = "", xaxt="n")
```

```

lines(soiltemp_T15_zoo, col = "grey70")
mtext("ARIMA with no seasonality", 1, -2.2, adj = 0.1, font = 2)
mtext(am0$call, 1, -1, adj = 0.1, col = 4, cex = 0.9,
      family="mono", font = 2)
mtext("(a)", 3, -1.2, cex = 1.2, font = 2, adj = 0.95)

plot(fcl, ylab = "Temperature (\u00B0C)", main = "", fcol = 6)
lines(soiltemp$conc, col = "grey70")
mtext("Time since start (seconds)", 1, 1.7, cex = 1.2, font = 2)
mtext(paste("ARIMA with", am1$arima[5], "h periodicity"),
      side = 1, line = -2.2, adj = 0.1, font = 2)
mtext(am1$call, 1, -1, adj = 0.1, col = 6, cex = 0.9, family="mono")
mtext("(b)", 3, -1.2, cex = 1.2, font = 2, adj = 0.95)

```

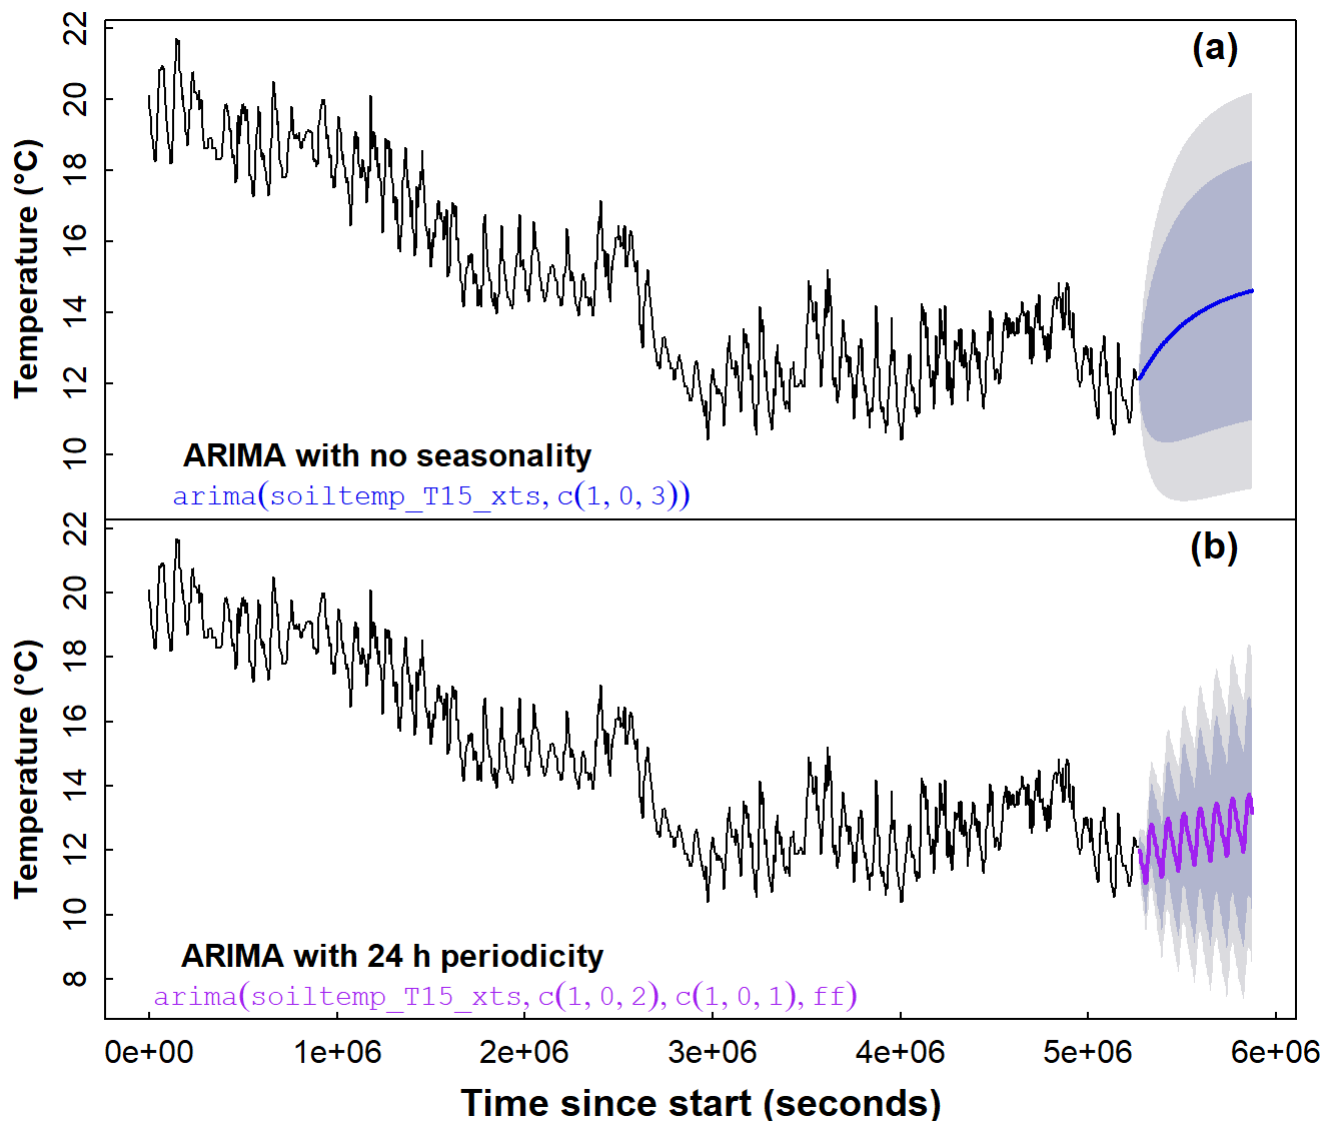


Figure 16: Time series forecasts for soil temperature, based on (a) a non-seasonal ARIMA model and (b) a seasonal ARIMA model.

From the two plots in Figure 16, we observe that the seasonal ARIMA model seems to reflect the form of the original data much better, by simulating the diurnal fluctuations in soil temperature. Although the forecasting uncertainty shown by the shaded area in Figure 16 is large for both types of ARIMA model, the interval seems smaller for the seasonal ARIMA model (Figure 16(b)).

We can also make a slightly 'prettier' time series forecast plot using the `autoplot()` function from the `ggplot2` package (Figure 17):

```

require(ggplot2) # gives best results using autoplot(...)

```

```
autoplot(fc1)+
  ylab("Temperature (\u00B0C)") +
  xlab(paste("Time since",index(soiltemp_T15_zoo)[1],"(s)")) +
  ggtitle("") +
  theme_bw()
```

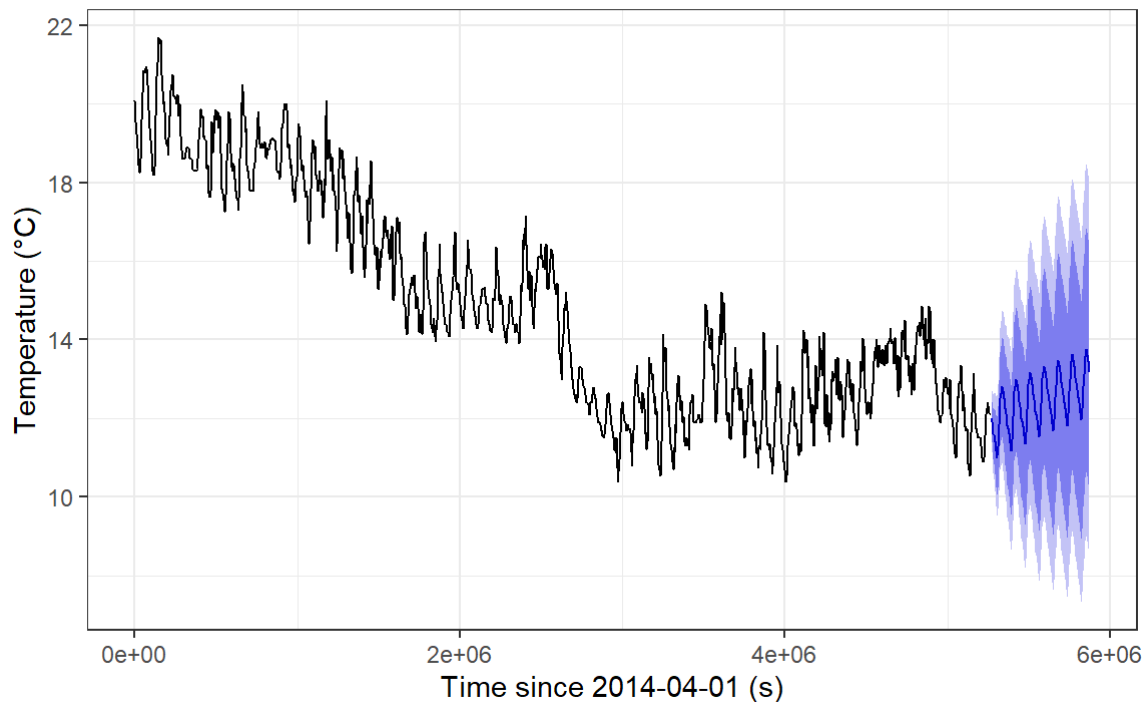


Figure 17: Time series forecast for soil temperature, based on a seasonal ARIMA model, plotted using the ggplot2 autoplot() function.

**ARIMA models are not the end of the time series modelling and forecasting story!**

## Exponential Smoothing Models

Sometimes, ARIMA models may not be the best option, and another commonly used method is **exponential smoothing**.

### Try an exponential smoothing model

Check `help(forecast::ets)` to correctly specify the model type using the `model =` option!

The easiest option (and the one used here) is to specify `model = "ZZZ"`, which chooses the type of model automatically.

```
soiltemp_T15_ets <- ets(soiltemp_T15_xts, model = "ZZZ")
summary(soiltemp_T15_ets)
## ETS (A,Ad,N)
##
## Call:
## ets(y = soiltemp_T15_xts, model = "ZZZ")
##
## Smoothing parameters:
##   alpha = 0.9999
##   beta  = 0.1083
##   phi   = 0.8
##
## Initial states:
##   l = 20.2683
##   b = -0.4487
##
## sigma: 0.3482
```

```
##
##      AIC      AICc      BIC
## 7588.660 7588.718 7620.394
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE
## Training set -0.003041526 0.3475639 0.2472475 -0.03856289 1.719534 0.01664266
##              ACF1
## Training set 0.0146761
```

In this case the `ets()` function has automatically selected a model ( $ETS(A,Ad,N)$ ) with additive trend and errors, but no seasonality. We don't do it here, but we could manually select a model with seasonality included.

The default plot function in R can plot forecasts from exponential smoothing models too (Figure 18).

```
plot(forecast(soiltemp_T15_ets, h=168), col=6,
     xlab = "Time since start (s)", ylab = "Temperature (\u00B0C)", main="")
mtext(names(soiltemp_T15_ets$par),3,seq(-1,-5,-1), adj = 0.7, col = 4, font = 3)
mtext(signif(soiltemp_T15_ets$par,3),3,seq(-1,-5,-1), adj = 0.8, col = 4, font = 3)
```

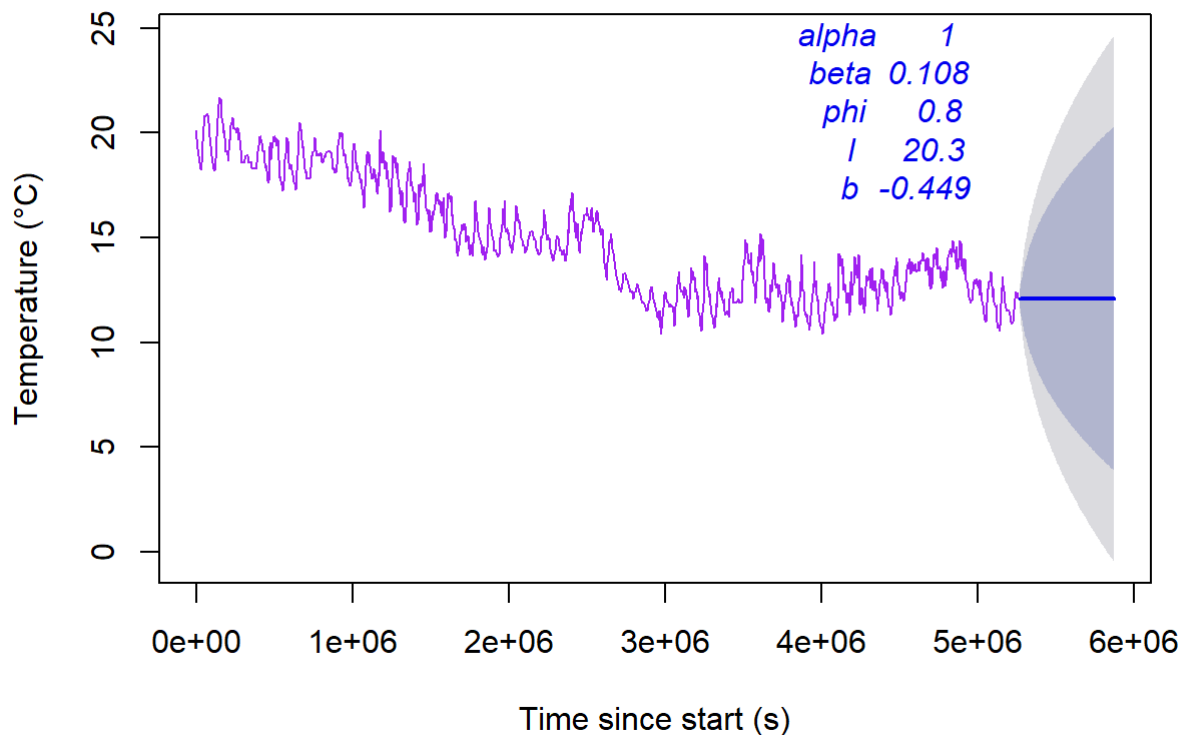


Figure 18: Forecast of soil temperature based on an automatically selected exponential smoothing model without a seasonal component.

Exponential smoothing also decomposes time series in a different way, into *level* and *slope* components.

```
plot(soiltemp_T15_ets, col = 2, xlab = "Time since start (s)") # ETS decomposition plots
```

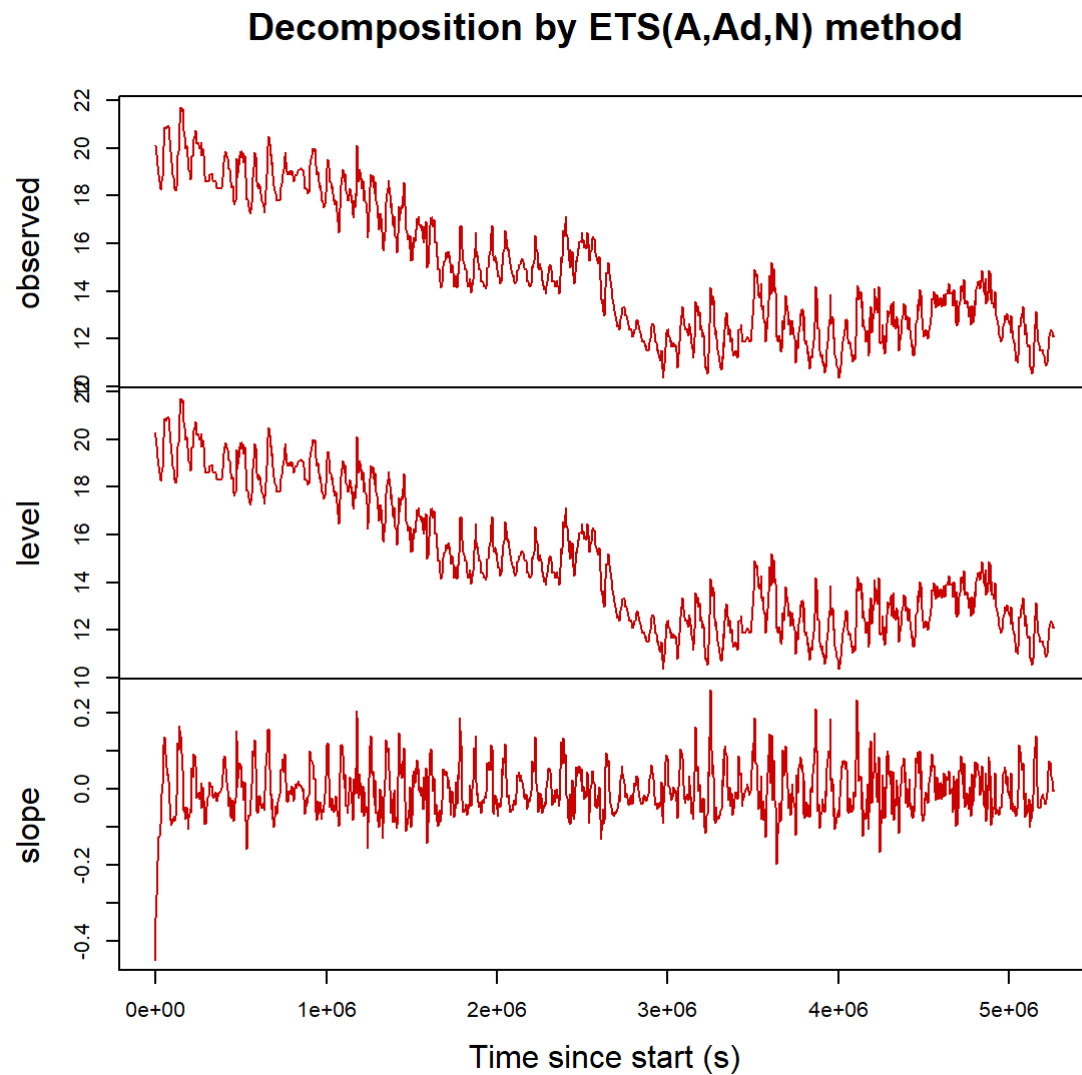


Figure 19: Exponential smoothing time series decomposition plots for soil temperature data.

*(Exponential smoothing is maybe not so good for the soil temperature data!)*

## References

- Department of Water (2015). *Calculating trends in nutrient data*, Government of Western Australia, Perth. [https://water.wa.gov.au/\\_\\_data/assets/pdf\\_file/0014/6800/Calculating-trends-in-nutrient-data-10-3-15.pdf](https://water.wa.gov.au/__data/assets/pdf_file/0014/6800/Calculating-trends-in-nutrient-data-10-3-15.pdf)
- Hyndman R, Athanasopoulos G, Bergmeir C, Caceres G, Chhay L, O'Hara-Wild M, Petropoulos F, Razbash S, Wang E, Yasmeen F (2020). *forecast: Forecasting functions for time series and linear models*. R package version 8.12, <http://pkg.robjhyndman.com/forecast>.
- Long, J.D., Teetor, P., 2019. Time series analysis. Chapter 14, *The R Cookbook*, Second Edition <https://rc2e.com/timeseriesanalysis>.
- R Core Team (2019). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing, Vienna, Austria. URL <https://www.R-project.org/>.
- Ryan, J.A. and Ulrich, J.M. (2018). *xts: eXtensible Time Series*. R package version 0.11-2. <https://CRAN.R-project.org/package=xts>
- Zeileis, A. and Grothendieck, G. (2005). zoo: S3 Infrastructure for Regular and Irregular Time Series. *Journal of Statistical Software*, **14**(6), 1-27. <https://doi.org/10.18637/jss.v014.i06>